

Image Formation

Forming an image

- First, we need some sort of sensor to receive and record light.
- Is this all we need?



- Do we get a useful image?

Restricting the Light

Pinhole Camera

object



barrier





film



- Advantages:
 - easy to simulate
 - everything is in focus
- Disadvantages:
 - needs a bright scene (or long exposure)
 - everything is in focus

Collecting the light

- Instead of throwing away all but a single ray, let's collect a bunch of rays and concentrate them at a single point on the sensor.
- To do this, we need to be able to change the path of a light ray.
- Fortunately, we have **refraction**. Light passing from one medium into denser one will bend towards the **normal** of the interface.




Stacking prisms



- We can use variously shaped prisms to take light rays of various angles and bend them to pass through a single point.
- As we use more and more prisms, the shape approaches a curve, and we get a **lens**.

Forming an image with a lens

- We can now replace the pinhole barrier with a lens, and we still get an image.



- Now there is a specific distance at which objects are “in focus”.
- By changing the shape of the lens, we change how it bends the light.

Optics

- **Focal point** - the point where parallel rays converge when passing through a lens.
- **Focal length** - the distance from the lens to the focal point.
- **Dioptr** - the reciprocal of the focal length, measured in meters.




Image Processing

Reading

The following links are available on the course homepage


- <http://www.dai.ed.ac.uk/HIPR2/noise.htm>
- <http://www.dai.ed.ac.uk/HIPR2/mean.htm>
- <http://www.cee.hw.ac.uk/hipr/html/median.html>

Definitions



- Many graphics techniques operate only on images
- **Image processing**: operations that take images as input, produce images as output
- In its most general form, an **image** is a function f from \mathbb{R}^2 to \mathbb{R}
 - $f(x, y)$ gives the intensity of a channel at position (x, y)
 - defined over a rectangle, with a finite range:
 - $f : [a,b] \times [c,d] \rightarrow [0,1]$
 - A color image is just three functions pasted together:
 - $\mathbf{f}(x, y) = (f_r(x, y), f_g(x, y), f_b(x, y))$

Images

- In computer graphics, we usually operate on **digital (discrete)** images
 - **Quantize** space into units (pixels)
 - Image is constant over each unit
 - A kind of step function
 - $f : \{0 \dots m-1\} \times \{0 \dots n-1\} \rightarrow [0,1]$
- An image processing operation typically defines an image f' in terms of an existing image f



Images as Functions




Pixel-to-pixel Operation

- The simplest operations are those that transform each pixel in isolation
- $f'(x, y) = g(f(x, y))$
- Example: threshold, RGB \longrightarrow greyscale

$f(x, y)$		$f'(x, y)$
.4 .3 .2 .1	$\xrightarrow{\text{Threshold } \geq 0.5}$.0 .0 .0 .0
.3 .8 .9 .2		.0 .8 .9 .0
.2 .5 .7 .3		.0 .5 .7 .0
.1 .0 .2 .4		.0 .0 .0 .0

Pixel Movement

- Some operations preserve intensities, but move around in the image
- $f'(x, y) = f(g(x, y), h(x, y))$
- Examples: many amusing warps of images



Noise

- Common types of noise



Original



Salt and pepper noise



Impulse noise



Gaussian noise

Noise

- Common types of noise:
 - **Salt and pepper noise:** random black and white pixels
 - **Impulse noise:** random white pixels
 - **Gaussian noise:** variations in intensity drawn from a Gaussian (normal) distribution

Noise Reduction

- Is there a way to “smooth” out the noise?

Reducing Noise

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

- Filtering
 - look at the neighborhood N around each pixel
 - replace each pixel with new value as a function of pixels in N
 - The behavior $g(i, j) = h(f, N)$ and f

Mean filtering

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$

$G[x, y]$

- Replace each pixel with an average of the pixels in the $k \times k$ box around it

– 3×3 case:
$$G[x, y] = \frac{1}{9} \sum_{u=0}^2 \sum_{v=0}^2 F[x + u - 1, y + v - 1]$$

Mean filtering

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$G[x, y]$

- Replace each pixel with an average of the pixels in the $k \times k$ box around it

– 3×3 case:
$$G[x, y] = \frac{1}{9} \sum_{u=0}^2 \sum_{v=0}^2 F[x + u - 1, y + v - 1]$$

What about border pixels?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$G[x, y]$

- Some options
 - don't evaluate—image gets smaller each time a filter is applied
 - pad the image with more rows and columns on the top, bottom, left, and right
 - option 1: copy the border pixels: add [0 0 0] to F in above case
 - option 2: reflect the image about the border: add [0 90 0] to F in above case

Effect of filter size

- What happens if we
- use a larger mean filter?
- 5×5? 7×7?

Weighted average filtering

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

1	1	1
1	1	1
1	1	1

$H[u, v]$

$F[x, y]$

- Replace each pixel with a weighted average of the pixels in the $k \times k$ box

– 3×3 case:
$$G[x, y] = \sum_{u=0}^2 \sum_{v=0}^2 H[u, v] * F[x + u - 1, y + v - 1]$$


Gaussian filter

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0


$F[x, y]$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} H[u, v]$$

- This filter H is a good approximation to $h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$
- Properties of Gaussian
 - more weight to the center
 - good model of blurring in optical systems
 - σ corresponds to width of the Gaussian



Comparison of mean vs. gaussian filter



Convolution

- Convolution is a fancy way to combine two functions.
 - Think of f as an image and g as a “smear” operator
 - g determines a new intensity at each point (pixel) in terms of intensities at the neighborhood of that point (pixel)

$$\begin{aligned}h(x, y) &= f(x, y) * g(x, y) \\&= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x', y') g(x - x', y - y') dx' dy'\end{aligned}$$

Discrete Convolution

- For digital images, integration becomes summation. We can express convolution as a two-dimensional sum:

$$\begin{aligned}h[i, j] &= f[i, j] * g[i, j] \\&= \sum_k \sum_l f[k, l] g[i - k, j - l]\end{aligned}$$

Convolution Representation

- Since f and g are defined over finite regions, we can write them out in two-dimensional arrays:

62	79	23	119	120	105	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30

X .2	X 0	X .2
X 0	X .2	X 0
X .2	X 0	X .2

Median Filter

- A **Median Filter** operates over a $k \times k$ region by selecting the median intensity in the region.
 - What advantage does a median filter have over a mean filter?
(answer available at “extra” - where you download the notes!)
 - Is a median filter a kind of convolution?

123	125	126	130	140
122	124	126	127	135
118	120	150	125	134
119	115	119	123	133
111	116	110	120	130

Neighbourhood values:

115, 119, 120, 123, 124,
125, 126, 127, 150

Median value: 124

Using Median Filters

Gaussian noise

Salt and pepper noise

3x3



5x5




7x7





Edge Detection

- One of the most important uses of image processing is **edge detection**
 - Really easy for humans
 - Really difficult for computers
 - Fundamental in computer vision
 - Important in many graphics applications
- What defines an edge?



Edge detection



- How can you tell that a pixel is on an edge?

Edge Detection

- Edge detection algorithms typically proceed in four steps:
 - Filtering: cut down on noise
 - Enhancement: amplify the difference between edges and non-edges
 - Thresholding
 - Localization (optional): estimate geometry of edges beyond pixels

Gradient

- The **gradient** is the 2D equivalent of the derivative:

$$\nabla f(x, y) = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right)$$




- Properties of the gradient
 - It's a vector
 - Points in the direction of maximum increase of f
(direction of the steepest descent)
 - Magnitude is rate of increase
- How can we approximate the gradient in a discrete image?

Image gradient

- The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The gradient points in the direction of most rapid change in intensity



The gradient direction is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- how does this relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Edge detection operator


- A popular gradient magnitude computation is the **Sobel operator**:

$$s_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$s_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

- We can then compute the magnitude of the vector $(G_x, G_y)^T$

Sobel Operator: Example



```

1 1 1 1 1 0 0 0 0 0
1 1 1 1 1 0 0 0 0 0
1 1 1 1 1 0 0 0 0 0
1 1 1 1 1 0 0 0 0 0
1 1 1 1 1 0 0 0 0 0
1 1 1 1 1 0 0 0 0 0
1 1 1 1 1 0 0 0 0 0

```

$$G_x = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} * \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} = -4$$

$$G_y = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = 0$$

$$(G_x, G_y)^T = (-4, 0)^T$$

Using Sobel Operators




Original



Smoothed



$G_x + 128^3$



$G_y + 128$



Magnitude





Threshold = 64



Threshold = 128


Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Where is the edge?

Solution: smooth first



Where is the edge? Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Summary

What you should take from this lecture:

- Formal definitions of image and image processing
- Kinds of image processing: pixel-to-pixel, movement, convolution, others
- Types of noise and strategies for noise reduction
- Definition of convolution and how discrete convolution works
- The effects of mean, median and Gaussian filtering
- How edge detection is done
- Gradients and discrete approximations