

Hidden Surfaces

Reading

Required

- Foley *et al*, Chapter 15

Optional

- Hearn and Baker, 13.1-13.3, 13.6-13.7

The Quest for 3D

Construct a 3D hierarchical geometric model

- Define a virtual camera
- Map points in 3D space to points in an image
- Produce a wireframe drawing in 2D from a 3D object
- Of course, there's more work to be done...

Introduction

- Not every part of every 3D object is visible to a particular viewer. We need an algorithm to determine what parts of each object should get drawn.
- Known as “hidden surface elimination” or “visible surface determination”.
- Hidden surface elimination algorithms can be in three major ways:
 - Object space vs. image space
 - Object order vs. image order
 - Sort first vs. sort last
 - Still a very active research area
- Where would we use a hidden surface algorithm?

Object Space Algorithms

- Operate on geometric primitives
 - For each object in the scene, compute the part of it which isn't occluded by any other object, then draw.
 - Must perform tests at high precision
 - Resulting information is resolution-independent
- Complexity
 - Must compare every pair of objects, so $O(n^2)$ for n objects
 - Optimizations can reduce this cost, but...
 - Best for scenes with few polygons or resolution-independent output
- Implementation
 - Difficult to implement!
 - Must carefully control numerical error

Image Space Algorithms


- Operate on pixels
 - For each pixel in the scene, find the object closest to the COP which intersects the projector through that pixel, then draw.
 - Perform tests at device resolution, result works only for that resolution
- Complexity
 - Must do something for every pixel in the scene, so at least $O(R)$.
 - Easiest solution is to test projector against every object, giving $O(nR)$.
 - More reasonable version only does work for pixels belonging to objects: $O(nr)$, r is number of pixels per object
 - Often, with more objects, each is smaller, we estimate $nr = O(R)$ in practice
- Implementation
 - Usually very simple!

Object Order vs. Image Order

- Object order
 - Consider each object only once - draw its pixels and move on to the next object
 - Might draw the same pixel multiple times
- Image order
 - Consider each pixel only once - draw part of an object and move on to the next pixel
 - Might compute relationships between objects multiple times

Sort First vs. Sort Last

- Sort first
 - Find some depth-based ordering of the objects relative to the camera, then draw from back to front
 - Build an ordered data structure to avoid duplicating work
- Sort last
 - Sort implicitly as more information becomes available




Important Algorithms

- Ray casting
- Z-buffer
- Binary space partitioning
- Back face culling

Ray Casting

- Partition the projection plane (PP) into pixels to match screen resolution
 - For each pixel $p_{i'}$ construct ray (projector) from COP through PP at that pixel and into scene
 - Intersect the ray with every object in the scene, color the pixel according to the object with the closest intersection




Ray Casting Implementation

- Parameterize the ray:

$$R(t) = (1-t)c + tp_i$$

- If a ray intersects some object O_i , get parameter t such that first intersection with O_i occurs at $R(t_i)$
- Which object owns the pixel?
- We will study ray-object intersections in more detail later (essential in ray tracing).



Aside: Definitions

- An algorithm exhibits *coherence* if it uses knowledge about the continuity of the objects on which it operates
- An *online* algorithm is one that doesn't need all the data to be present when it starts running
 - Example: insertion sort

Ray Casting: Analysis

Categorization:

- Easy to implement?
- Hardware implementation?
- Coherence?
- Memory intensive?
- Pre-processing required?
- Online?
- Handles transparency?
- Handles refraction?
- Polygon-based?
- Extra work for moving objects?
- Extra work for moving viewer?
- Efficient shading?
- Handles cycles and self-intersections?


Z-buffer

- Idea: along with a pixel's red, green and blue values, maintain some notion of its *depth*
 - An additional channel in memory, like alpha
 - Called the depth buffer or Z-buffer
- When the time comes to draw a pixel, compare its depth with the depth of what's already in the framebuffer. Replace only if it's closer
 - Very widely used (in hardware, e.g. GeForce3/4)
- History
 - Originally described as “brute-force image space algorithm”
 - Written off as impractical algorithm for huge memories
 - Today, done easily in hardware

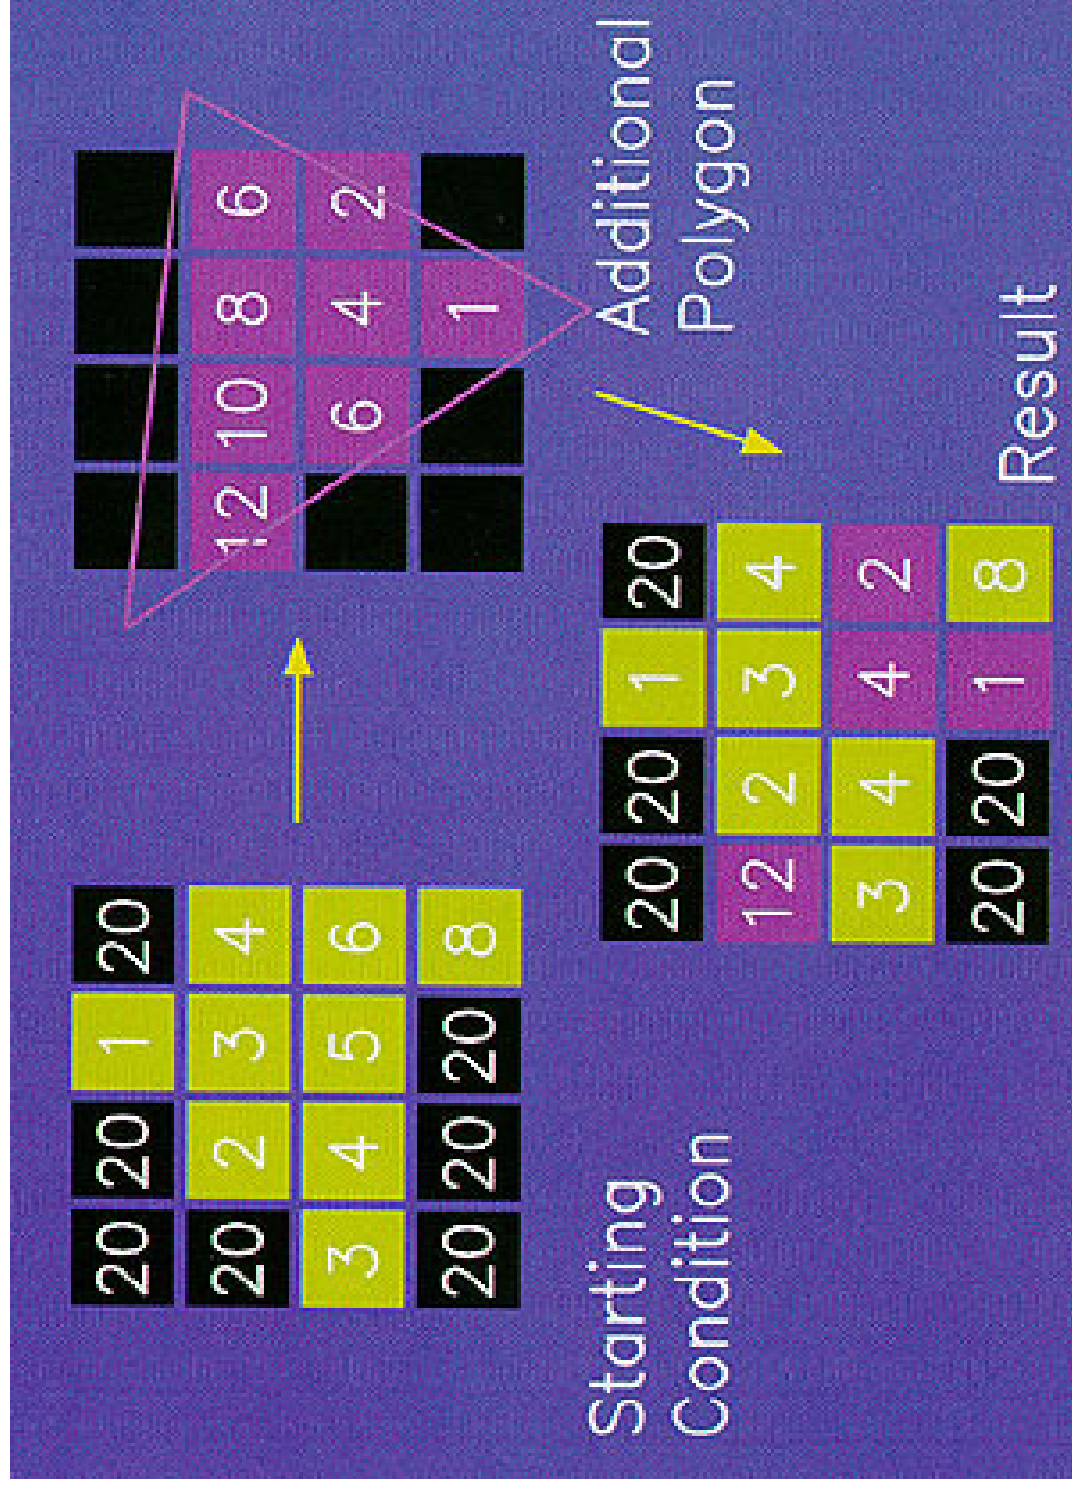
Z-buffer Implementation

```
for each pixel pi {  
    Z-buffer[ pi ] = FAR  
    Fb[ pi ] = BACKGROUND_COLOR  
}  
for each polygon P {  
    for each pixel pi in the projection of P {  
        Compute depth z and shade s of P at pi  
        if z < Z-buffer[ pi ] {  
            Z-buffer[ pi ] = z  
            Fb[ pi ] = s  
        }  
    }  
}
```

Z-Buffer Algorithm




Z-Buffer Algorithm




Z-buffer Tricks

- The shade of a triangle can be computed incrementally from the shades of its vertices (taking advantage of coherence)
- Can do the same with depth



Z Value Interpretation




$$Z_a = Z_1 - (Z_1 - Z_2) \frac{y_1 - y_s}{y_1 - y_2}$$

$$Z_b = Z_1 - (Z_1 - Z_3) \frac{y_1 - y_s}{y_1 - y_3}$$

$$Z_p = Z_b - (Z_b - Z_a) \frac{x_b - x_p}{x_b - x_a}$$

Depth Preserving Conversion to Parallel Projection



Z Buffer: Analysis

Categorization:





- Easy to implement?
- Hardware implementation?
- Coherence?
- Memory intensive?
- Pre-processing required?
- Online?
- Handles transparency?
- Handles refraction?
- Polygon-based?
- Extra work for moving objects?
- Extra work for moving viewer?
- Efficient shading?
- Handles cycles and self-intersections?





Binary Space Partitioning

- Goal: build a tree that captures some relative depth information between objects. Use it to draw objects in the right order.
 - Tree doesn't depend on camera position, so we can change viewpoint and redraw quickly
 - Called the binary space partitioning tree, or BSP tree
- Key observation: The polygons in the scene are painted in the correct order if for each polygon P ,
 - Polygons on the far side of P are painted first
 - P is painted next
 - Polygons in front of P are painted last

Binary Space-Partitioning Tree

- Given a polygon p
- Two lists of polygons:
 - those that are behind(p) : B
 - those that are in front(p) : F
- If eye is in-front(p), right display order is B, p, F
- Otherwise it is F, p, B





BSP Tree Construction

```
BSPtree makeBSP( L: list of polygons ) {  
    if L is empty {  
        return the empty tree  
    }  
    Choose a polygon P from L to serve as root  
    Split all polygons in L according to P  
    return new TreeNode(  
        P,  
        makeBSP( polygons on negative side of P ),  
        makeBSP( polygons on positive side of P ) )  
}
```

- Splitting polygons is expensive! It helps to choose P wisely at each step.
- Example: choose five candidates, keep the one that splits the fewest polygons

BSP Tree Display

```
showBSP( v: Viewer, T: BSPtree ) {  
    if T is empty then return  
    P := root of T  
    if viewer is in front of P {  
        showBSP( back subtree of T )  
        draw P  
        showBSP( front subtree of T )  
    } else {  
        showBSP( front subtree of T )  
        draw P  
        showBSP( back subtree of T )  
    }  
}
```

BSP Tree: Analysis

Categorization:


- Easy to implement?
- Hardware implementation?
- Coherence?
- Memory intensive?
- Pre-processing required?
- Online?
- Handles transparency?
- Handles refraction?
- Polygon-based?
- Extra work for moving objects?
- Extra work for moving viewer?
- Efficient shading?
- Handles cycles and self-intersections?

Back Face Culling


- Can be used in conjunction with polygon-based algorithms
 - Often, we don't want to draw polygons that face away from the viewer. So test for this and eliminate (cull) backfacing polygons before drawing
- How can we test for this?

Back Face Culling

- What is a normal vector?
- It is a vector pointing in an orthogonal direction of a plane.




- A polygon is front facing when θ , the angle between the viewer and the normal to the polygon, lies between -90° to $+90^\circ$ or $\cos \theta \geq 0$



Back Face Culling

- Use dot product to test for visibility.
- Visibility test: $N_p \cdot V > 0$
where N_p is the polygon normal
 V is the line of sight (viewing) vector
- How to calculate polygon normal?



- A polygon is front facing when θ , the angle between the viewer and the normal to the polygon, lies between -90° to $+90^\circ$ or $\cos \theta \geq 0$
- All vertices of polygons must be listed in the same order (either clockwise or anticlockwise)

Back Face Culling

- Eye space is the most convenient space in which to 'cull' polygon
- Remove all polygons that face away from the viewer
- For a scene that consists of only a single close object, it solves the hidden surface problem *completely*.
- Not work for scene with multiple objects
- Cannot remove all unnecessary polygons of a concave object
- In most cases, it is a preprocess to reduce invisible polygons.
- But it cannot remove all unnecessary triangles.

Summary

- Classification of hidden surface algorithms
- Understanding ray casting algorithms
- Understanding of Z-buffer
- Familiarity with BSP trees and back face culling